Linear Regression

Felipe José Bravo Márquez

September 9, 2021

Introduction
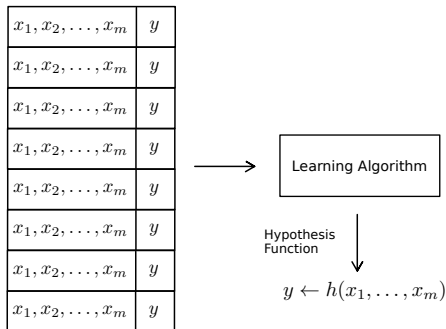
- A regression model is used to model the relationship of a numerical dependent variable **y** with m independent variables $\mathbf{x}_1, \mathbf{x}_2, \ldots, \mathbf{x}_m$ [Wasserman, 2013].
- The dependent variable **y** is also called **target**, **outcome**, or **response** variable.
- The independent variables $\mathbf{x}_1, \mathbf{x}_2, \ldots, \mathbf{x}_m$ are also called **covariates**, **attributes**, **features**, or **predictor variables**.
- Roughly speaking we want to know the expected value of **y** from the values of **x**:

$$\mathbb{E}(y|x_1, x_2, \ldots, x_m)$$

- We use these models when we believe that the response variable **y** can be modeled by other independent variables.
- To perform this type of analysis we need a dataset consisting of n observations that include both the response variable and each of the attributes.
- We refer to the process of **fitting** a regression function as the process in which we infer a hypothesis function h from the data that allows us to **predict** unknown values for **y** using the values of the attributes.

Introduction (2)

- This process of fitting a function from data is referred to in the areas of data mining and machine learning as **training**.
- In those disciplines, functions are said to **learn** from data.
- Since we need observations where the value of **y** (the output) is known to learn the function, such techniques are referred to as **supervised learning**.
- When **y** is a categorical variable we have a **classification** problem.

$x_1, x_2, \ldots, x_m$	y
$x_1, x_2, \ldots, x_m$	y
$x_1, x_2, \ldots, x_m$	y
$x_1, x_2, \ldots, x_m$	y
$x_1, x_2, \ldots, x_m$	y
$x_1, x_2, \ldots, x_m$	y
$x_1, x_2, \ldots, x_m$	y
$x_1, x_2, \ldots, x_m$	y

$\longrightarrow$ Learning Algorithm

Hypothesis Function

$$y \leftarrow h(x_1, \ldots, x_m)$$

Simple Linear Regression

- In simple linear regression, we have a single independent variable **x** to model the dependent variable **y**.
- The following linear relationship between the variables is assumed:

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i \quad \forall i$$

- The parameter β_0 represents the intercept of the line (the value of y when x is zero).
- The parameter β_1 is the slope and represents the change of **y** when we vary the value of **x**. The greater the magnitude of this parameter the greater the linear relationship between the variables.
- The ϵ_i values correspond to the errors or **residuals** associated with the model.
- We have to find a linear function or straight line h_β that allows us to find an estimate of y, $\hat{y}$ for any value of x with the minimum expected error.

$$h(x) = \beta_0 + \beta_1 x$$

Simple Linear Regression

Figure: Source: https://www.jmp.com

Least Squares

- The ordinary least squares method allows us to estimate $\hat{\beta}_0$ and $\hat{\beta}_1$ by minimizing the sum of squared errors (SSE) of the observed data.
- Suppose we have n observations of **y** and **x**, we compute the sum of squared errors (SSE) as follows:

$$SSE = \sum_{i=1}^{n}(y_i - h(x_i))^2 = \sum_{i=1}^{n}(y_i - \beta_0 - \beta_1 x_i)^2 \tag{1}$$

- To find the parameters that minimize the error we calculate the partial derivatives of SSE with respect to β_0 and β_1. Then we equal the derivatives to zero and solve the equation to find the parameter values.

$$\frac{\partial SSE}{\partial \beta_0} = -2\sum_{i=1}^{n}(y_i - \beta_0 - \beta_1 x_i) = 0 \tag{2}$$

$$\frac{\partial SSE}{\partial \beta_1} = -2\sum_{i=1}^{n}(y_i - \beta_0 - \beta_1 x_i)x_i = 0 \tag{3}$$

Least Squares

- From the above system of equations the normal solutions are obtained:

$$\hat{\beta}_1 = \frac{\sum_i^n (x_i - \overline{x})(y_i - \overline{y})}{\sum_i^n (x_i - \overline{x})^2} \tag{4}$$

$$\hat{\beta}_0 = \overline{y} - \hat{\beta}_1 \overline{x} \tag{5}$$

- The fitted model represents the line of least squared error.

Coefficient of Determination R^2

- How can we evaluate the quality of our fitted model?
- A very common metric is the coefficient of determination R^2.
- It is calculated from errors that are different from SSE.
- The total sum of squares (SST) is defined as the predictive error when we use the mean $\overline{y}$ to predict the response variable y (it is very similar to the variance of the variable):

$$\text{SST} = \sum_{i}^{n} (y_i - \overline{y})^2$$

- The regression sum of squares (SSM), on the other hand, indicates the variability of the model's predictions with respect to the mean:

$$\text{SSM} = \sum_{i}^{n} (\hat{y}_i - \overline{y})^2$$

Coefficient of Determination R^2

- It can be proved that all the above errors are related as follows:

$$SST = SSE + SSM \tag{6}$$

- The coefficient of determination for a linear model R^2 is defined as:

$$R^2 = \frac{\text{SSM}}{\text{SST}} = \frac{\sum_i^n (\hat{y}_i - \overline{y})^2}{\sum_i^n (y_i - \overline{y})^2} \tag{7}$$

- An alternative but equivalent definition:

$$R^2 = 1 - \frac{\text{SSE}}{\text{SST}} = 1 - \frac{\sum_i^n (y_i - \hat{y}_i)^2}{\sum_i^n (y_i - \overline{y})^2} \tag{8}$$

- Another alternative but equivalent definition:

$$R^2 = 1 - \frac{\text{var}(\epsilon)}{\text{var}(y)} \tag{9}$$

- It is often interpreted as the "variance explained" by the model.
- The coefficient usually takes values between 0 to 1 and the closer its value is to 1 the higher the quality of the model[1].
- The value of R^2 is equivalent to the squared linear correlation (Pearsons) between y and $\hat{y}$.

$$R^2 = \text{cor}(y, \hat{y})^2$$

---

[1] It can take negative values when predictions are worse than using $\overline{y}$ in all predictions.

Assumptions of the Linear Model

Whenever we fit a linear model we are implicitly making certain assumptions about the data.

Assumptions

1. Linearity: the response variable is linearly related to the attributes.
2. Normality: errors have a zero mean normal distribution: $\epsilon_i \sim N(0, \sigma^2)$.
3. Homoscedasticity: errors have constant variance (same value of σ^2).
4. Independence: errors are independent of each other.

Probabilistic Interpretation

- Considering the above assumptions, we are saying that the probability density (PDF) of the errors ϵ is a Gaussian with zero mean and constant variance σ^2:

$$\text{PDF}(\epsilon_i) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{\epsilon_i^2}{2\sigma^2}\right)$$

- Moreover all ϵ_i are IID (independent and identically distributed).
- This implies that:

$$\text{PDF}(y_i|x_i; \beta) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - h_\beta(x_i))^2}{2\sigma^2}\right)$$

- Which implies that the distribution of **y** given the values of **x**, parameterized by β is also normally distributed.
- Let's estimate the parameters of the model (β) using maximum likelihood estimation.

Probabilistic Interpretation

- The likelihood function of β can be written as:

$$\mathcal{L}(\beta) = \prod_{i=1}^{n} \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - h_\beta(x_i))^2}{2\sigma^2}\right)$$

- and the the log likelihood $l_n(\beta)$:

$$l_n(\beta) = \log \mathcal{L}(\beta) \tag{10}$$

$$= \log \prod_{i=1}^{n} \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - h_\beta(x_i))^2}{2\sigma^2}\right) \tag{11}$$

$$= \sum_{i=1}^{n} \log \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - h_\beta(x_i))^2}{2\sigma^2}\right) \tag{12}$$

$$= n \log\left(\frac{1}{\sqrt{2\pi}\sigma}\right) - \frac{1}{\sigma^2} \times \frac{1}{2} \sum_{i=1}^{n} (y_i - h_\beta(x_i))^2 \tag{13}$$

Probabilistic Interpretation

- After ignoring the constants, we realize that maximizing $l_n(\beta)$ is equivalent to minimizing:

$$\sum_{i=1}^{n}(y_i - h_\beta(x_i))^2$$

- Which is equivalent to minimizing SSE!
- Then, if one estimates the parameters of β using maximum likelihood estimation one arrives at the same results as doing least squares estimation.
- Therefore, every time we perform least squares estimation, we are implicitly making the same probabilistic assumptions mentioned above (e.g., normality, homoscedasticity).

Standard errors for regression models

- If we want to make inferences about the regression parameter estimates, then we also need an estimate of their variability [Poldrack, 2019].
- To compute this, we first need to compute the **mean squared error** MS_{error} of the model:

$$MS_{error} = \frac{SSE}{df} = \frac{\sum_{i=1}^{n}(y_i - h(x_i))^2}{n - p}$$

- The degrees of freedom (df) are determined by subtracting the number of estimated parameters (2 in this case: β_0 and β_1) from the number of observations (n).
- We can then compute the standard error for the model as:

$$SE_{model} = \sqrt{MS_{error}}$$

A significance test for β

- In order to get the standard error for a specific regression parameter estimate, SE_{β_x}, we need to rescale the standard error of the model by the square root of the sum of squares of the X variable:

$$SE_{\hat{\beta}_x} = \frac{SE_{model}}{\sqrt{\sum_{i=1}^{n}(x_i - \overline{x})^2}}$$

- Now we can compute a t statistic to tell us the likelihood of the observed parameter estimates compared to some expected value under the null hypothesis.

- In this case we will test against the null hypothesis of no effect (i.e. $\beta_{H_0} = 0$):

$$t_{n-p} = \frac{\hat{\beta} - \beta_{H_0}}{SE_{\hat{\beta}_x}} = \frac{\hat{\beta}}{SE_{\hat{\beta}_x}}$$

- Later we will see that R automatically reports the t-statistics and p-values of all coefficients of a linear model.

- This allows us to determine whether the linear relationship between the two variables (y and x) is significant.

Example: a model of height

- We are going to work with the dataset `Howell1` that has demographic data from Kalahari !Kung San people collected by Nancy Howell in the late 1960s.
- The !Kung San are the most famous foraging population of the twentieth century, largely because of detailed quantitative studies by people like Howell. [McElreath, 2020]

Figure: By Staehler - Own work, CC BY-SA 4.0, https://commons.wikimedia.org/w/index.php?curid=45076017

Example: a model of height

- Each observation corresponds to an individual.

- The variables of the dataset are:
    1. height: height in cm
    2. weight: weight in kg
    3. age: age in years
    4. male: gender indicator (1 for male 0 for woman)

- Let's explore the linear correlations between the variables

```
> d <- read.csv("howell1.csv")
> cor(d)
          height    weight        age        male
height 1.0000000 0.9408222 0.683688567 0.139229021
weight 0.9408222 1.0000000 0.678335313 0.155442866
age    0.6836886 0.6783353 1.000000000 0.005887126
male   0.1392290 0.1554429 0.005887126 1.000000000
```

Example: a model of height

- We can see that there is a positive correlation between `height` and `age`.
- Let's filter out the non-adult examples because we know that height is strongly correlated with age before adulthood.

  ```
  d2 <- d[ d$age >= 18 , ]
  ```

- Now age doesn't correlate with height:

  ```
  > cor(d2)
               height      weight         age       male
  height   1.0000000   0.7547479  -0.10183776  0.69999340
  weight   0.7547479   1.0000000  -0.17290430  0.52445271
  age     -0.1018378  -0.1729043   1.00000000  0.02845498
  male     0.6999934   0.5244527   0.02845498  1.00000000
  ```

- Let's model height as a function of weight using a simple linear regression:

$$\text{height(weight)} = \beta_0 + \beta_1 * \text{weight}$$

Example: a model of height

- In R the linear models are created with the command `lm` that receives formulas of the form `y~x` ($y = f(x)$).

```
> reg1<-lm(height~weight,d2)
> reg1

Call:
lm(formula = height ~ weight, data = d2)

Coefficients:
(Intercept)      weight
    113.879       0.905
```

Example: a model of height

- We can see that the coefficients of the model are $\beta_0 = 113.879$ and $\beta_1 = 0.905$.
- The estimate of $\beta_0 = 113.879$, indicates that a person of weight 0 should be around 114cm tall, which we know is nonsense but it what our linear model believes.
- Since β_1 is a slope, the value 0.905 can be read as a person 1 kg heavier is expected to be 0.90 cm taller.
- We can directly access the coefficients and store them in a variable:

```
> reg1.coef<-reg1$coefficients
> reg1.coef
(Intercept)      weight
113.8793936   0.9050291
```

Example: a model of height

- We can view various indicators about the linear model with the command **summary**:

```
> summary(reg1)

Call:
lm(formula = height ~ weight, data = d2)

Residuals:
     Min       1Q   Median       3Q      Max
-19.7464  -2.8835   0.0222   3.1424  14.7744

Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) 113.87939    1.91107   59.59   <2e-16 ***
weight        0.90503    0.04205   21.52   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.086 on 350 degrees of freedom
Multiple R-squared:  0.5696,	Adjusted R-squared:  0.5684
F-statistic: 463.3 on 1 and 350 DF,  p-value: < 2.2e-16
```

- We can see that β_0 and β_1 are both statistically significantly different from zero.

Example: a model of height

- We see that the coefficient of determination R^2 has a value of 0.57 which is not so good but acceptable.
- We can conclude that the weight while providing useful information to model a part of the variability of the height of the !Kung people, is not enough to build a highly reliable model.
- We can store the results of the command `summary` in a variable then access the coefficient of determination:

```
> sum.reg1<-summary(reg1)
> sum.reg1$r.squared
[1] 0.5696444
```

- We can also access the fitted values which are the values predicted by my model for the data used:

```
> reg1$fitted.values
        1        2        3        4        5        6
157.1630 146.9001 142.7180 161.8839 151.2362 170.8895
```

Example: a model of height

- We can check now that all definitions given for R^2 are equivalent.

```
> SSE<-sum(reg1$residuals^2)
> SST<-sum((d2$height-mean(d2$height))^2)
> SSM<-sum((reg1$fitted.values-mean(d2$height))^2)
> SSM/SST
[1] 0.5696444
> 1-SSE/SST
[1] 0.5696444
> 1-var(reg1$residuals)/var(d2$height)
[1] 0.5696444
> cor(d2$height,reg1$fitted.values)^2
[1] 0.5696444
```

- Or manually calculate the standard error for the model SE_{model}:

```
> df <- dim(d2)[1]-2
> SE <- sqrt(SSE/df)
> SE
[1] 5.086336
```

Example: a model of height

- Suppose now that we know the weight for two !Kung people but we don't know their height.
- We could use our linear model to predict the height of these two people.
- To do this in R we must use the command `predict.lm` which receives the linear model and a data.frame with the new data:

```
> new.weights<-data.frame(weight=c(50,62))
> predict.lm(object=reg1,newdata=new.weights)
       1         2
159.1308 169.9912
> # this is equivalent to:
> reg1.coef[1]+reg1.coef[2]*new.weights[1:2,]
[1] 159.1308 169.99122
```

Multivariate Linear Regression

- Suppose we have m independent variables: $x_1, x_2, \ldots, x_m$.
- In many cases more variables can better explain the variability of the response variable **y** than a single one.
- A multivariate linear model is defined as follows:

$$y_i = \beta_0 + \beta_1 x_{i,1} + +\beta_2 x_{i,2} + \cdots + \beta_m x_{i,m} + \epsilon_i \quad \forall i \in \{1, n\}$$

- In essence we are adding a parameter for each independent variable.
- Then, we multiply the parameter by the variable and add that term to the linear model.
- In the multivariate model, all the properties of the simple linear model are extended.
- The problem can be represented in a matrix form:

$$Y = X\beta + \epsilon$$

Multivariate Linear Regression

- Where Y is a vector $n \times 1$ response variables:

$$Y = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix}$$

- X is a $n \times (m+1)$ matrix with the explanatory variables. We have n observations of the m variables. The first column is constant equal to 1 ($x_{i,0} = 1 \quad \forall i$) to model the intercept variables β_0.

$$X = \begin{pmatrix} x_{1,0} & x_{1,1} & x_{1,2} & \cdots & x_{1,m} \\ x_{2,0} & x_{2,1} & x_{2,2} & \cdots & x_{2,m} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_{n,0} & x_{n,1} & x_{n,2} & \cdots & x_{n,m} \end{pmatrix}$$

Multivariate Linear Regression

- Then, β is a $(m+1) \times 1$ vector of parameters.

$$\beta = \begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_m \end{pmatrix}$$

- Finally, ϵ is a $n \times 1$ vector with the model errors.

$$\epsilon = \begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{pmatrix}$$

- Using matrix notation, we can see that the sum of squared errors (SSE) can be expressed as:

$$\text{SSE} = (Y - X\beta)^T(Y - X\beta)$$

- Minimizing this expression by deriving the error as a function of β and setting it equal to zero leads to the normal equations:

$$\hat{\beta} = (X^TX)^{-1}X^TY$$

Multivariate Linear Regression in R

- Now we will study a multiple linear regression for the Howell1 data.
- Let's add the variable **age** as an additional predictor for **height**.
- We know that age is good at predicting height for non-adults so we will work with the original dataset.
- Let's fit the following linear multi-variate model:

$$\text{height(weight,age)} = \beta_0 + \beta_1 * \text{weight} + \beta_2 * \text{age}$$

- In R, we add more variables to the linear model with the operator **+** :

```
reg2<-lm(height˜weight+age,d)
```

Multivariate Linear Regression in R

```
> summary(reg2)

Call:
lm(formula = height ~ weight + age, data = d)

Residuals:
    Min       1Q   Median       3Q      Max
-29.0350  -5.4228   0.7333   6.4919  19.6964

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 75.96329    1.04216  72.890  < 2e-16 ***
weight       1.65710    0.03656  45.324  < 2e-16 ***
age          0.11212    0.02594   4.322 1.84e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 9.214 on 541 degrees of freedom
Multiple R-squared:  0.889,Adjusted R-squared:  0.8886
F-statistic:  2166 on 2 and 541 DF,  p-value: < 2.2e-16
```

Multivariate Linear Regression in R

- When we had a simple regression we could see the fitted model as a line.
- Now that we have two independent variables we can see the fitted model as a plane.
- If we had more independent variables our model would be a hyper-plane.
- We can plot the plane of our linear model of two independent variables and one dependent variable in R as follows:

```
library("scatterplot3d")
s3d <- scatterplot3d(d[,c("weight","age","height")],
                     type="h", highlight.3d=TRUE,
                     angle=55, scale.y=0.7, pch=16,
                     main="height~weight+age")
s3d$plane3d(reg2, lty.box = "solid")
```

height~weight+age

R^2 adjusted

- The coefficient of determination R^2 tends increases when extra explanatory variables are added to the model.
- This happens even when the additional variables are useless.
- This is not desirable a property when we want to compare different models.
- R^2 adjusted or $\overline{R}^2$ penalizes for the number of variables m.

$$\overline{R}^2 = 1 - (1 - R^2)\frac{n-1}{n-m-1}$$

where n is the number of examples.

Polynomial Regression

- Polynomial regression uses powers of a variable (e.g., squares, cubes) as extra attributes.
- The most common polynomial regression is a parabolic model of the mean:

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \epsilon_i \quad \forall i$$

- Let's fit a polynomial regression for the height variable using a parabolic model for weight
- Because the square of a large number can be truly massive we are going to standardize the weight by substracting the mean and dividing by the standard deviation.

```
d$weight_s <-(d$weight-mean(d$weight))/sd(d$weight)
```

- Then we fit the model as follows:

```
> reg4 <- lm(height~weight_s+I(weight_s^2),d)
> reg4

Call:
lm(formula = height ~ weight_s + I(weight_s^2), data = d)

Coefficients:
  (Intercept)        weight_s   I(weight_s^2)
      146.660          21.415          -8.412
```

Categorical Predictors

- We can also add categorical predictors to our model.
- For example, let's try using the sex of the Kalahari people to predict height:

$$\text{height(male)} = \beta_0 + \beta_1 * \text{male}$$

- Here "male" is a binary or dummy variable (it takes the value 1 when the person is male and 0 otherwise).
- In R, it is important to make sure that our categorical variables are represented as factors:

```
> d$male<-as.factor(d$male)
> reg5<-lm(height~male,d)
```

- The coefficients of the model are:

```
> reg5$coefficients
(Intercept)        male1
 134.630278     7.690759
```

Categorical Predictors

- Here β_0 is the average height of females and β_1 is the average height difference between males and females:

```
> sum(reg5$coefficients)
[1] 142.321
> means<-tapply(d$height,d$male,mean)
> means
       0        1
134.6303 142.3210
> means[2]-means[1]
       1
7.690759
```

```
> summary(reg5)

Call:
lm(formula = height ~ male, data = d)

Residuals:
   Min     1Q Median     3Q    Max
-81.87 -12.73  12.65  18.33  36.75

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  134.630      1.615  83.366  < 2e-16 ***
male1          7.691      2.350   3.273  0.00113 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 27.36 on 542 degrees of freedom
Multiple R-squared:  0.01938,Adjusted R-squared:  0.01758
F-statistic: 10.71 on 1 and 542 DF,  p-value: 0.001131
```

Categorical Predictors

- The p-value of the statistical sinfificance test for β_1 in reg5 is the same as that obtained from an unpaired two-sample t-tests with equal variance:

```
> t.test(d$height˜d$male, var.equal=T)

Two Sample t-test

data:  d$height by d$male
t = -3.2733, df = 542, p-value = 0.001131
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -12.306146  -3.075372
sample estimates:
mean in group 0 mean in group 1
       134.6303        142.3210
```

- Notice that we must set `var.equal=T` to get the same results, since the linear model assumes that the variance is constant.

Categorical Predictors

- We can also fit a multivariate linear model with both numeric and categorical variables:

$$\text{height(weight,male)} = \beta_0 + \beta_1 * \text{weight} + \beta_2 * \text{male}$$

```
> reg6<-lm(height˜weight+male,d)
> reg6

Call:
lm(formula = height ˜ weight + male, data = d)

Coefficients:
(Intercept)      weight         male1
    75.5489      1.7664       -0.3971
```

- In this model we use the same slope β_1 relating height to weight for both groups.
- The coefficient β_2 indicates an expected weight difference between males and females once the "weight" is known.

Interactions

- In many cases we want to fit a model using lines with separate slopes for each category.
- This is done using an **interaction**, which in our case is just an additional parameter associated with the the product between the category and the predictor.

$$\text{height}_{int}(\text{weight,male}) = \beta_0^{(i)} + \beta_1^{(i)} * \text{weight} + \beta_2^{(i)} * \text{male} + \beta_3^{(i)} * \text{weight} * \text{male}$$

- In R, interactions are denoted by using the * or : symbol in the formula.

```
> reg7<-lm(height~weight+male+weight:male,d)
> # or lm(height~weight*male,d)
> reg7

Call:
lm(formula = height ~ weight + male + weight:male, data = d)

Coefficients:
 (Intercept)       weight         male1  weight:male1
    74.2536       1.8051        2.0683       -0.0695
```

Interactions

- The above model can be interpreted as a compact version of fitting two models, one for each group.
- A model for the male group would be as follows:

$$\text{height}_{male}(\text{weight}_{male}) = \beta_0^{(m)} + \beta_1^{(m)} * \text{weight}_{male}$$

```
> d.male<-d[d$male==1,]
> reg8<-lm(height~weight,d.male)
Call:
lm(formula = height ~ weight, data = d.male)
Coefficients:
(Intercept)       weight
     76.322        1.736
```

and for the female group:

$$\text{height}_{female}(\text{weight}_{male}) = \beta_0^{(f)} + \beta_1^{(f)} * \text{weight}_{male}$$

```
> d.female<-d[d$male==0,]
> reg9<-lm(height~weight,d.female)
Call:
lm(formula = height ~ weight, data = d.female)
Coefficients:
(Intercept)       weight
     74.254        1.805
```

Interactions

- We can find many relationships between the coefficients of the model with interaction and the two independent models.
- All of them can be deduced by cancelling out terms in the formula of height$_{int}$ for the female cases (*male* = 0).
- First, the intercept of the model with the interaction height$_{int}$ is the same as the intercept of the model for females $\beta_0^{(i)} = \beta_0^{(f)}$:

```
> reg7$coefficients["(Intercept)"]
(Intercept)
   74.25359
> reg9$coefficient["(Intercept)"]
(Intercept)
   74.25359
```

- The same applies for the slope of weight for height$_{int}$ and height$_{female}$ $\beta_1^{(i)} = \beta_1^{(f)}$:

```
> reg7$coefficients["weight"]
  weight
1.805118
> reg9$coefficient["weight"]
  weight
1.805118
```

Interactions

- The interaction slope from height$_{int}$ encodes the difference in weight slopes between both groups $\beta_3^{(i)} = \beta_1^{(m)} - \beta_1^{(f)}$:

```
> reg7$coefficients["weight:male1"]
weight:male1
 -0.06949648
> reg8$coefficient["weight"]-reg9$coefficient["weight"]
     weight
-0.06949648
```

- Finally, the category coefficient $\beta_2^{(i)}$ from height$_{int}$ corresponds to the difference of intercepts between height$_{male}$ and height$_{female}$ $\beta_2^{(i)} = \beta_0^{(m)} - \beta_0^{(f)}$:

```
> reg7$coefficients["male1"]
    male1
2.068278
> reg8$coefficients["(Intercept)"]
-reg9$coefficient["(Intercept)"]
(Intercept)
    2.068278
```

- We can use a simple linear model to describe the relation between two variables and to decide whether that relationship is statistically significant.
- In addition, the model allows us to predict the value of the dependent variable given some new value(s) of the independent variable(s).
- Most importantly, a mutivariate linear model allow us to build models that incorporate multiple independent variables.

Bonus: Four Cardinal Rules of Statistics by Daniela Witten

Now that we have concluded the chapter on Frequentist inference, it is good to discuss the points raised by Daniela Witten in a tweet.

Dr. Daniela Witten @daniela_witten · Oct 2, 2020
FOUR CARDINAL RULES OF STATISTICS 📈 📊 📉

A thread

1/🚪

💬 80   🔁 3K   ❤️ 9.4K   ↥

One: Correlation does not imply causation

- Yes, I know you know this, but it's so easy to forget! Yeah, YOU OVER THERE, you with the p-value of 0.0000001 — yes, YOU!! That's not causation.
- No matter how small the p-value for a regression of IQ onto shoe size is, that doesn't mean that big feet cause smarts!! It just means that grown-ups tend to have bigger feet and higher IQs than kids.
- So, unless you can design your study to uncover causation (very hard to do in most practical settings — the field of causal inference is devoted to understanding the settings in which it is possible), the best you can do is to discover correlations. Sad but true.

Bonus: Four Cardinal Rules of Statistics by Daniela Witten

Two: a p-value is just a test of sample size

- Read that again — I mean what I said! If your null hypothesis doesn't hold (and null hypotheses never hold IRL) then the larger your sample size, the smaller your p-value will tend to be.

- If you're testing whether mean=0 and actually the truth is that mean=0.000000001, and if you have a large enough sample size, then YOU WILL GET A TINY P-VALUE.

- Why does this matter? In many contemporary settings (think: the internet), sample sizes are so huge that we can get TINY p-values even when the deviation from the null hypothesis is negligible. In other words, we can have STATISTICAL significance w/o PRACTICAL significance.

- Often, people focus on that tiny p-value, and the fact that the effect is of **literally no practical relevance** is totally lost.

- This also means that with a large enough sample size we can reject basically ANY null hypothesis (since the null hypothesis never exactly holds IRL, but it might be "close enough" that the violation of the null hypothesis is not important).

- Want to write a paper saying Lucky Charms consumption is correlated w/blood type? W/a large enough sample size, you can get a small p-value. (Provided there's some super convoluted mechanism with some teeny effect size. . . which there probably is, b/c IRL null never holds)

Bonus: Four Cardinal Rules of Statistics by Daniela Witten

Three: seek and you shall find

- If you look at your data for long enough, you will find something interesting, even if only by chance!
- In principle, we know that we need to perform a correction for multiple testing if we conduct a bunch of tests.
- But in practice, what if we decide what test(s) to conduct AFTER we look at data? Our p-value will be misleadingly small because we peeked at the data. Pre-specifying our analysis plan in advance keeps us honest... but in reality, it's hard to do!!!

- Everyone is asking me about the mysterious and much-anticipated fourth rule of statistics. The answer is simple: we haven't figured it out yet.... that's the reason we need to do research in statistics

McElreath, R. (2020).
Statistical rethinking: A Bayesian course with examples in R and Stan.
CRC press.

Poldrack, R. A. (2019).
Statistical thinking for the 21st century.
https://statsthinking21.org/.

Wasserman, L. (2013).
All of statistics: a concise course in statistical inference.
Springer Science & Business Media.