

Directed Graphical Models

Felipe José Bravo Márquez

August 30, 2021

Directed Graphical Models

- Probabilistic graphical models (PGMs) are a rich framework for encoding probability distributions over complex domains [Koller and Friedman, 2009].
- In this class we will focus on directed graphical models (DGMs), which are one type of PGM.
- Directed graphical models (DGMs) are a family of probability distributions that admit a compact parametrization that can be naturally described using a **directed graph**.
- DGMs are also known as **Bayesian networks**.
- We won't use that term in this class because statistical inference for DGMs can be performed using frequentist or Bayesian methods [Wasserman, 2013].
- These types of graphs are also called **causal graphs** or **causal diagrams**.

Directed Graphical Models

- DGMs link two very different branches of mathematics: probability and graph theory.
- They also have intriguing connections to philosophy, particularly the question of **causality**.
- At the same time, they are widely used in statistics and machine learning.
- DGMs can be used to solve problems in fields as diverse as medicine, language processing, vision, and many others [Ermon and Kuleshov,].
- But before introducing them more formally we need to learn the following mathematical concepts:
 - 1 Conditional independence
 - 2 The chain rule of probability
 - 3 Directed acyclical graphs

Conditional Independence

- Two random variables X and Y are independent, written $X \perp Y$ if $f(x, y) = f(x)f(y)$ for all values of x, y .
- This also implies that $f(x|y) = f(x)$ and $f(y|x) = f(y)$.
- Notice that f can be either a density function for continuous random variables or a probability mass function for discrete random variables.
- In the same way $f(x|y)$ and $f(y|x)$ correspond to conditional densities or mass functions.
- Now, suppose we have three random variables A, B, C .
- A and B are conditionally independent given C , written $A \perp B|C$, if:

$$f(a|b, c) = f(a|c)$$

for all a, b and c .

Conditional Independence

- Intuitively, $A \perp B|C$ means that, once you know C , B provides no extra information about A .
- Notice that $A \perp B|C$ doesn't necessarily imply that $A \perp B$.

Example

- Let H, V, A be three random variables representing a person's height, vocabulary and age.
- H and V are dependent $f(v|h) \neq f(v)$ since very small people tend to be children, known for their more basic vocabularies.
- But knowing that two people are 19 years old (i.e., conditional on age) there is no reason to think that one person's vocabulary is larger if we are told that they are taller:

$$f(v|h, a) = f(v|a) \Leftrightarrow V \perp H|A$$

The chain rule of probability

- For a set of random variables X_1, \dots, X_n , the chain rule of probability allow us to express the joint probability function $f(x_1, x_2, \dots, x_n)$ as a product of n conditional probabilities:

$$f(x_1, x_2, \dots, x_n) = f(x_1)f(x_2|x_1)f(x_3|x_2, x_1) \dots f(x_n|x_{n-1}, \dots, x_2, x_1).$$

- For example for the case of $n = 3$

$$f(x_1, x_2, x_3) = f(x_1)f(x_2|x_1)f(x_3|x_2, x_1)$$

using the definition of conditional probabilities we have that

$$f(x_2|x_1) = f(x_1, x_2)/f(x_1)$$

and

$$f(x_3|x_2, x_1) = f(x_1, x_2, x_3)/f(x_1, x_2)$$

- By replacing these expressions into the chain of products, many expressions cancel out and we obtain $f(x_1, x_2, x_3)$.

Joint Probabilities

- Expressing a joint probability function can be expensive.
- For example if there are m binary random variables, the complete distribution is specified by $2^m - 1$ joint probabilities.
- For example, if we have two Boolean variables (A, B) , we need the probabilities $\mathbb{P}(A, B)$, $\mathbb{P}(\neg A, B)$, $\mathbb{P}(A, \neg B)$, and $\mathbb{P}(\neg A, \neg B)$.
- A joint distribution for a set of random variables gives all the information there is about the distribution.
- Note that for m boolean variables, the joint distribution contains 2^m values.
- However, the sum of all the joint probabilities must be 1 because the probability of all possible outcomes must be 1.
- Thus, to specify the joint distribution, one needs to specify 2^{m-1} numbers.

Joint Probabilities

- The main idea of DGMs is that by assuming that some variables are conditionally independent we can use the probability chain rule to represent the joint distribution more compactly.
- For example in the previous example, a conditional independence assumption could be: $X_3 \perp X_2 | X_1: f(x_3 | x_2, x_1) = f(x_3 | x_2)$.
- Then our joint distribution would be reduced to:

$$f(x_1, x_2, x_3) = f(x_1)f(x_2 | x_1)f(x_3 | x_2)$$

- In the following slides we will try to explain these ideas with an example taken from [Koller and Friedman, 2009].

Example

- Consider the problem faced by a company trying to hire a recent college graduate.
- The company's goal is to hire intelligent employees, but there is no way to test intelligence directly.
- However, the company has access to the student's SAT scores¹, which are informative but not fully indicative.
- Thus, our probability space is induced by the two random variables Intelligence (I) and SAT (S).
- For simplicity, we assume that each of these takes two values: $Val(I) = \{i_1, i_0\}$, which represent the values high intelligence (i_1) and low intelligence (i_0).
- Similarly $Val(S) = \{s_1, s_0\}$, which also represent the values high (score) and low (score), respectively.

¹The SAT is a standardized test widely used for college admissions in the United States.

Example

- Thus, our joint distribution in this case has four entries.
- For example, one possible joint distribution f would be:

I	S	$f(I, S)$
i^0	s^0	0.665
i^0	s^1	0.035
i^1	s^0	0.06
i^1	s^1	0.24

- Using the chain rule of conditional probabilities, we have that

$$f(I, S) = f(I)f(S|I)$$

- Intuitively, we are representing the process in a way that is more compatible with causality.
- Various factors (e.g., genetics, upbringing) first determined (stochastically) the student's intelligence.
- Her performance on the SAT is determined (stochastically) by his intelligence.
- We note that the models we construct are not required to follow causal intuitions, but they often do.

Directed Acyclic Graphs (DAGs)

- A directed graph consists of a set of nodes with arrows between some nodes.
- Graphs are useful for representing independence relations between variables.
- More formally, a directed graph G consists of a set of vertices V and an edge set E of ordered pairs of vertices.
- For our purposes, each vertex corresponds to a random variable.
- If $(Y, X) \in E$ then there is an arrow pointing from Y to X .




Figure: A directed graph with vertices $V = \{X, Y, Z\}$ and edges $E = \{(Y, X), (Y, Z)\}$.

Directed Acyclic Graphs (DAGs)

- If an arrow connects two variables X and Y (in either direction) we say that X and Y are adjacent.
- If there is an arrow from X to Y then X is a parent of Y and Y is a child of X .
- The set of all parents of X is denoted by π_X or $\pi(X)$.
- A directed path between two variables is a set of arrows all pointing in the same direction linking one variable to the other such as the chain shown below:





Figure: A chain graph with a directed path.

- X is an ancestor of Y if there is a directed path from X to Y (or $X = Y$).
- We also say that Y is a descendant of X .


Directed Acyclic Graphs (DAGs)

- A directed path that starts and ends at the same variable is called a cycle.
- A directed graph is acyclic if it has no cycles.
- In this case we say that the graph is a directed acyclic graph or DAG.

Acyclic Graph



Cyclic Graph



- From now on, we only deal with directed acyclic graphs since it is very difficult to provide a coherent probability semantics over graphs with directed cycles.

Probability and DAGs


- For the remainder of this class we will use the terms directed graphical model (DGM) and directed acyclical graph (DAG) interchangeably.
- A DAG is a distribution in which each factor on the right hand side depends only on a small number of ancestor variables $\pi(x)$. [Ermon and Kuleshov,]
- Let G be a DAG with vertices $V = (X_1, \dots, X_d)$.
- If P is a distribution for V with probability function $f(x)$ (density or mass), we say that G represents P , if

$$f(x) = \prod_{j=1}^d f(x_j | \pi_{x_j})$$

where π_{x_j} is the set of parent nodes of X_j

Probability and DAGs

- The next figure shows a DAG with four variables.



- The probability function takes the following decomposition:
- $f(\text{overweight}, \text{smoking}, \text{heart disease}, \text{cough}) = f(\text{overweight}) \times f(\text{smoking}) \times f(\text{heart, disease} | \text{overweight}, \text{smoking}) \times f(\text{cough} | \text{smoking})$.

- The interpretation of direct acyclic graphs as carriers of independence assumptions does not necessarily imply causation.
- In fact, it will be valid for any set of recursive independencies along any ordering of the variables, not necessarily causal or chronological.
- However, the ubiquity of DAG models in statistical and AI applications stems (often unwittingly) primarily from their causal interpretation
- That is, as a system of processes, one per family, that could account for the generation of the observed data.
- It is this causal interpretation that explains why DAG models are rarely used in any variable ordering other than those which respect the direction of time and causation. [Pearl, 2009]


An Example

- The best way to understand DAGs is to imagine trying to model a situation in which causality plays a role.
- And also our understanding of what is actually going on is incomplete
- So we need to describe things probabilistically.
- The following example is based on [Charniak, 1991].
- Eugene Charniak is a famous AI researcher who's got the following situation.
- When he goes home at night, he wants to know if his family is home before trying the doors.
- Often when his wife leaves the house, she turns on an outdoor light.

An Example

- Eugene's wife can also turn on the outdoor light if she is expecting a guest.
- Also, they have a female dog.
- When nobody is home, the dog is put in the back yard.
- The same is true if the dog has bowel troubles.
- Finally, if the dog is in the backyard, Eugene's will probably hear her barking
- The next slide shows a DAG encoding all the above causal relationships.

An Example



- The DAG can help to predict what will happen in a particular scenario (if his family goes out, the dog goes out)
- Or to infer causes from observed effects (if the light is on and the dog is out, then his family is probably out).

- sdsad

- sdsad

Estimation for DAGs

- Two estimation questions arise in the context of DAGs.
- First, given a DAG \mathcal{G} and data d_1, \dots, d_n from a distribution f consistent with \mathcal{G} , how do we estimate f ?
- Second, given data d_1, \dots, d_n how do we estimate \mathcal{G} ?
- The first question is pure estimation while the second involves model selection.
- These are very involved topics and are beyond the scope of this course.
- We will just briefly mention the main ideas.

Estimation for DAGs

- If we are doing frequentist inference, we typically use some parametric model $f(x|\pi_x; \theta_x)$ for each conditional density.
- The likelihood function is then

$$\mathcal{L}(\theta) = \prod_{i=1}^n f(d_i; \theta) = \prod_{i=1}^n \prod_{j=1}^m f(X_{ij}|\pi_j; \theta_j)$$

- where X_{ij} is the value of X_j for the i th data point and j are the parameters for the j th conditional density.
- We can then estimate the parameters by maximum likelihood.
- On the other hand, if we want to perform Bayesian inference we must set priors for all our variables X_1, \dots, X_m and estimate the posterior accordingly.

Estimation for DAGs

- To estimate the structure of the DAG itself, we could fit every possible DAG using maximum likelihood and use AIC (or some other method) to choose a DAG.
- However, there are many possible DAGs so we would need much data for such a method to be reliable.
- Also, searching through all possible DAGs is a serious computational challenge.
- Producing a valid, accurate confidence set for the DAG structure would require astronomical sample sizes.
- If prior information is available about part of the DAG structure, the computational and statistical problems are at least partly ameliorated [Wasserman, 2013].

Conclusions

- Blabla

References I



Charniak, E. (1991).
Bayesian networks without tears.
AI magazine, 12(4):50–50.



Ermon, S. and Kuleshov, V.
Cs228 notes.



Koller, D. and Friedman, N. (2009).
Probabilistic graphical models: principles and techniques.
MIT press.



Pearl, J. (2009).
Causality.
Cambridge university press.



Wasserman, L. (2013).
All of statistics: a concise course in statistical inference.
Springer Science & Business Media.