

# Introduction to Statistical Inference

Felipe José Bravo Márquez

April 16, 2021

# Populations and Samples

- The main goal of statistical inference is investigate properties about a target **population**.
- Example: What is the average height of the Chilean people?
- In order to draw conclusions about a **population**, it is generally not feasible to gather all the data about it.
- We must make reasonable conclusions about a population based on the evidence provided by **sample data**.
- A **sample staticic** or simply **statistic** is a quantitative measure calculated from the data. Examples: the mean, the standard deviation, the minimum, the maximum.
- Our goal in sampling is to determine the value of a statistic for an entire population of interest, using just a small subset of the population.
- We do this primarily to save time and effort.
- Idea: Why go to the trouble of measuring every individual in the population when just a small sample is sufficient to accurately estimate the statistic of interest? [Poldrack, 2019]

# Samples and Surveys

- Random samples
- Stratified samples
- Biases

# Statistical Inference

- The process of drawing conclusions about a population from sample data is known as **statistical inference**.
- From a general point of view, the goal of inference is to **infer** the distribution that generates the observed data.
- Example: Given a sample  $X_1, \dots, X_n \sim F$ , how do we infer  $F$ ?
- However, in most cases we are only interested in inferring some property of  $F$  (e.g., its **mean** value).
- Statistical models that assume that the distribution can be modeled with a finite set of parameters  $\theta = (\theta_1, \theta_2, \dots, \theta_k)$  are called **parametric models**.
- Example: if we assume that the data comes from a normal distribution  $N(\mu, \sigma^2)$ ,  $\mu$  and  $\sigma$  would be the parameters of the model.

# Frequentist Approaches

The statistical methods to be presented in this class are known as **frequentist (or classical)** methods. They are based on the following postulates [Wasserman, 2013]:

- Probability refers to limiting relative frequencies. Probabilities are objective properties of the real world.
- Parameters are fixed, unknown constants. Because they are not fluctuating, no useful probability statements can be made about parameters.
- Statistical procedures should be designed to have well-defined long run frequency properties. For example, a 95 percent confidence interval should trap the true value of the parameter with limiting frequency at least 95 percent.

There is another approach to inference called **Bayesian inference**, which is based on different postulates, to be discussed later in the course.

# Point Estimation

- Point estimation is the process of finding the **best guess** for some quantity of interest from a **statistical sample**.
- In a general sence, this quantity of interest could be a parameter in a parametric model, a CDF  $F$ , a probability density function  $f$ , a regression function  $r$ , or a prediction for a future value  $Y$  of some random variable.
- In this class we will consider this quantity of interest as a **population parameter**  $\theta$ .
- By convention, we denote a point estimate of  $\theta$  by  $\hat{\theta}$  or  $\hat{\theta}_n$ .
- It is important to remark that while  $\theta$  is an unknown fixed value,  $\hat{\theta}$  depends on the data and is therefore a random variable.
- We need to bear in mind that the process of sampling is by definition a random experiment.

# Point Estimation

## Formal Definition

- Let  $X_1, \dots, X_n$  be  $n$  IID data points from some distribution  $F$ .
- A point estimator  $\hat{\theta}_n$  of a parameter  $\theta$  is some function of  $X_1, \dots, X_n$ :

$$\hat{\theta}_n = g(X_1, \dots, X_n)$$

- The **bias** of an estimator is defined as:

$$\text{bias}(\hat{\theta}_n) = \mathbb{E}(\hat{\theta}_n) - \theta$$

- An estimator is unbiased if  $\mathbb{E}(\hat{\theta}_n) = \theta$  or  $\text{bias}(\hat{\theta}_n) = 0$

# Sampling Distribution

- If we take multiple samples, the value of our statistical estimate  $\hat{\theta}_n$  will also vary from sample to sample.
- We refer to this distribution of our estimator across samples as the **sampling distribution** [Poldrack, 2019].
- The sampling distribution may be considered as the distribution of  $\hat{\theta}_n$  for all possible samples from the same population of size  $n^1$ .
- The sampling distribution describes the variability of the point estimate around the true population parameter from sample to sample.
- We need to bear in mind this is an imaginary concept, since we can't obtain all possible samples.

---

<sup>1</sup><https://courses.lumenlearning.com/boundless-statistics/chapter/sampling-distributions/>



# Standard Error

- The standard deviation of  $\hat{\theta}_n$  is called the **standard error**  $se$ :

$$se(\hat{\theta}_n) = \sqrt{\mathbb{V}(\hat{\theta}_n)}$$

- The standard error tells us about the variability of the estimator between all possible samples of the same size.
- It is a measure of the uncertainty of the point estimate.

# The Sample Mean

- Let  $X_1, X_2, \dots, X_n$  be a random sample of a population of mean  $\mu$  and variance  $\sigma^2$
- A sample statistic we can derive from the data is the **sample mean**  $\overline{X}_n$

$$\overline{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$$

- The sample mean is a statistical estimate of the mean  $\overline{X}_n = \hat{\mu}$ .
- The sampling distribution of the sample mean:
- We can show that the sample mean is an unbiased estimator of  $\mu$ :

$$\mathbb{E}(\overline{X}_n) = \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n} \times \mathbb{E}\left(\sum_{i=1}^n X_i\right) = \frac{1}{n}(n \times \mu) = \mu$$

# Point Estimation (4)

- The standard error of the sample mean  $se(\bar{X}_n) = \sqrt{\mathbb{V}(\bar{X}_n)}$  can be calculated as:

$$\mathbb{V}(\bar{X}_n) = \mathbb{V}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n^2} \mathbb{V}\left(\sum_{i=1}^n X_i\right) = \frac{n}{n^2} \mathbb{V}(X_i) = \frac{\sigma^2}{n}$$

- Then,  $se(\bar{X}_n) = \frac{\sigma}{\sqrt{n}}$

# Point Estimation (5)

- We usually do not know  $\sigma$  of the population.
- When we want to estimate the variance of a population from a sample we are talking about the **sample variance**:
- There are two common estimators, a biased version

$$s_n^2 = \frac{1}{n} \sum_i^n (X_i - \bar{X}_n)^2$$

- An unbiased version

$$s^2 = \frac{1}{n-1} \sum_i^n (X_i - \bar{X}_n)^2$$

- When we do not know the variance of the population and want to estimate the mean, the standard error is estimated:

$$\hat{se}(\bar{X}_n) = \frac{s}{\sqrt{n}}$$

# Point Estimation (6)

- Let  $X_1, \dots, X_n \sim \text{Bernoulli}(p)$  and let  $\hat{p}_n = \frac{1}{n} \sum_i X_i$
- Then  $\mathbb{E}(\hat{p}_n) = \frac{1}{n} \sum_i \mathbb{E}(X_i) = p$ , and  $\hat{p}_n$  is unbiased.
- The standard error  $se$  would be

$$se = \sqrt{\mathbb{V}(\hat{p}_n)} = \sqrt{p(1-p)/n}$$

- The estimated standard error  $\hat{se}$ :

$$\hat{se} = \sqrt{\hat{p}(1-\hat{p})/n}$$

# Point Estimation (7)

- A good estimator is expected to be unbiased and of minimum variance.
- Unbiasedness used to receive much attention but these days is considered less important
- Many of the estimators we will use are biased.
- A reasonable requirement for an estimator is that it should converge to the true parameter value as we collect more and more data.
- A point estimator  $\hat{\theta}_n$  of a parameter  $\theta$  is **consistent** if it converges to the true value when the number of data in the sample tends to infinity..
- The quality of an estimator can be measured using the **mean squared error** (MSE)

$$MSE = \mathbb{E}_{\theta}(\hat{\theta}_n - \theta)^2$$

# Point Estimation (8)

- If for an estimator  $\hat{\theta}_n$ , its  $bias \rightarrow 0$  and its  $se \rightarrow 0$  when  $n \rightarrow \infty$ ,  $\hat{\theta}_n$  is a consistent estimator of  $\theta$ .
- For example, for the sample mean  $\mathbb{E}(\overline{X}_n) = \mu$  which implies that the  $bias = 0$  and  $se(\overline{X}_n) = \frac{\sigma}{\sqrt{n}}$  tends to zero when  $n \rightarrow \infty$ . Then  $\overline{X}_n$  is a consistent estimator of the mean.
- For the case of the Bernoulli experiment one has that  $\mathbb{E}(\hat{p}) = p \Rightarrow bias = 0$  and  $se = \sqrt{p(1-p)/n} \rightarrow 0$  when  $n \rightarrow \infty$ . Then  $\hat{p}$  is a consistent estimator of  $p$ .

# Maximum Likelihood Estimation



# Confidence Interval

- We know that the value of a point estimator **varies** from sample to sample.
- It is more reasonable to find an interval that is likely to trap the real value of the parameter with a certain probability.
- The general form of a confidence interval in the following:

$$\text{Confidence Interval} = \text{Sample Statistic} \pm \text{Margin Error}$$

- The wider the interval the more uncertainty there is about the value of the parameter.

# Confidence Interval

## Definition

- A **confidence interval** for an unknown population parameter  $\theta$  with a **confidence level**  $1 - \alpha$ , is an interval  $C_n = (a, b)$  where:

$$\mathbb{P}(\theta \in C_n) = 1 - \alpha$$

- In addition  $a = a(X_1, \dots, X_n)$  and  $b = b(X_1, \dots, X_n)$  are functions of the data.
- The  $\alpha$  value is known as the **significance** level, generally taken as 0.05, which is equivalent to working with a confidence level of 95%.
- Significance can be interpreted as the probability of being wrong.

# Confidence Interval

- There is a lot of **confusion** about how to interpret a confidence interval.
- A confidence interval is not a probability statement about  $\theta$  since  $\theta$  is a fixed quantity in Frequentist inference setting, not a random variable
- One way to interpret them is to say that if we repeat the **same experiment** many times, the interval will contain the value of the parameter  $(1 - \alpha)\%$  of the times.
- This interpretation is correct, but we rarely repeat the same experiment several times.
- A better interpretation: one day I collect data I create a 95% confidence interval for a parameter  $\theta_1$ . Then on day 2, I do the same for a parameter  $\theta_2$  and so repeatedly  $n$  times. The 95% of my intervals will contain the actual values of the parameters.

# Confidence Interval

- Later in the course, we will discuss Bayesian methods in which we treat  $\theta$  as if it were a random variable and we do make probability statements about  $\theta$ .
- In particular, we will make statements like “the probability that  $\theta$  is in  $C_n$ , given the data, is 95 percent.”
- However, these Bayesian intervals refer to degree-of-belief probabilities.
- These Bayesian intervals will not, in general, trap the parameter 95 percent of the time.

# Confidence Interval of the Mean

- We have  $n$  independent observations  $X_1, \dots, X_n$  (IID) of distribution  $N(\mu, \sigma^2)$ .
- Suppose  $\mu$  is **unknown** but  $\sigma^2$  is **known**.
- We know that  $\overline{X}_n$  is an unbiased estimator of  $\mu$ .
- By the law of large numbers we know that the distribution of  $\overline{X}_n$  is concentrated around  $\mu$  when  $n$  is large.
- By the CLT we know that

$$Z = \frac{\overline{X}_n - \mu}{\frac{\sigma}{\sqrt{n}}} \sim N(0, 1)$$

when  $n$  is large.


# Confidence Interval

- We want to find an interval  $C_n = (\mu_1, \mu_2)$  with confidence level  $1 - \alpha$ :


$$\mathbb{P}(\mu_1 \leq \mu \leq \mu_2) = 1 - \alpha$$

- Let  $z_a = \Phi^{-1}(1 - a)$ , with  $a \in [0, 1]$  where  $\Phi^{-1}$  is the quantile function of a standardized normal.
- This is equivalent to saying that  $z_a$  is the value such that  $1 - \Phi(z_a) = \mathbb{P}(Z \geq z_a) = a$ .
- By symmetry of the normal distribution:  $z_{\alpha/2} = -z_{(1-\alpha/2)}$ .

# Confidence Interval




# Confidence Interval





# Confidence Interval



# Confidence Interval

- The confidence interval for  $\mu$  is:

$$C_n = (\overline{X}_n - z_{\alpha/2} \frac{\sigma}{\sqrt{n}}, \overline{X}_n + z_{\alpha/2} \frac{\sigma}{\sqrt{n}})$$

- Then  $z_{\alpha/2}$  tells us how many times we have to multiply the **standard error** to build the interval.
- The smaller the value of  $\alpha$  the larger the value of  $z_{\alpha/2}$  and hence the wider the interval.
- Proof:

$$\begin{aligned} \mathbb{P}(\mu \in C_n) &= \mathbb{P}(\overline{X}_n - z_{\alpha/2} \frac{\sigma}{\sqrt{n}} < \mu < \overline{X}_n + z_{\alpha/2} \frac{\sigma}{\sqrt{n}}) \\ &= \mathbb{P}(-z_{\alpha/2} < \frac{\overline{X}_n - \mu}{\frac{\sigma}{\sqrt{n}}} < z_{\alpha/2}) \\ &= \mathbb{P}(-z_{\alpha/2} < Z < z_{\alpha/2}) \\ &= 1 - \alpha \end{aligned}$$

# Confidence Interval

- Since  $z_{\alpha/2} = \Phi^{-1}(1 - \alpha/2)$  we can use the quantile function of the normal to calculate confidence intervals in R.

```
> alpha <- 0.05
> xbar <- 5
> sigma <- 2
> n <- 20
> se <- sigma/sqrt(n)
> error <- qnorm(1-alpha/2)*se
> left <- xbar-error
> right <- xbar+error
> left
[1] 4.123477
> right
[1] 5.876523
>
```

# T Distribution

- In practice, if we do not know  $\mu$  we are unlikely to know  $\sigma$ .
- If we estimate  $\sigma$  using  $s$ , confidence intervals are better build using the distribution **T-student**, especially when the sample size is small.

## T Distribution


- An R.V. has distribution  $t$  with  $k$  degrees of freedom when it has the following PDF:

$$f(t) = \frac{\Gamma(\frac{k+1}{2})}{\sqrt{k\pi}\Gamma(\frac{k}{2})(1 + \frac{t^2}{k})^{(k+1)/2}}$$

- When  $k = 1$  it is called **Cauchy** distribution.
- When  $k \rightarrow \infty$  it converges to a standardized normal distribution.
- The t-distribution has wider tails than the normal distribution when it has few degrees of freedom.

# T Distribution

```
x<-seq(-8,8,length=400)
y1<-dnorm(x)
y2<-dt(x=x,df=1)
y3<-dt(x=x,df=10)
y4<-dt(x=x,df=350)
plot(y1~x,type="l",col="green")
lines(y2~x,type="l",col="blue")
lines(y3~x,type="l",col="black")
lines(y4~x,type="l",col="red")
```



# T-Distribution Confidence Interval

- Let  $s^2 = \frac{1}{n-1} \sum_i^n (X_i - \bar{X}_n)^2$  we have:

$$T = \frac{\bar{X}_n - \mu}{\frac{s}{\sqrt{n}}} \sim t_{n-1}$$

- Let  $t_{n-1,a} = \mathbb{P}(T > a)$ , equivalent to the quantile function  $qt$  evaluated at  $(1 - a)$ .
- The resulting confidence interval is:

$$C_n = (\bar{X}_n - t_{n-1,\alpha/2} \frac{s}{\sqrt{n}}, \bar{X}_n + t_{n-1,\alpha/2} \frac{s}{\sqrt{n}})$$

- Since the tails of the  $t$  distribution are wider when  $n$  is small, the resulting confidence intervals are wider.

# T-Distribution Confidence Interval

- Let's calculate a confidence interval for the mean of `Petal.Length` of the **Iris** data with 95% confidence.

```
>data(iris)
>alpha<-0.05
>n<-length(iris$Petal.Length)
>xbar<-mean(iris$Petal.Length)
>xbar
[1] 3.758
>s<-sd(iris$Petal.Length)
>se<-s/sqrt(n)
>error<-qt(p=1-alpha/2,df=n-1)*se
>left<-xbar-error
>left
[1] 3.473185
>right<-xbar+error
>right
[1] 4.042815
```

- Otra forma:

```
>test<-t.test(iris$Petal.Length,conf.level=0.95)
>test$conf.int
[1] 3.473185 4.042815
```

# The Bootstrap



# References I



Poldrack, R. A. (2019).  
*Statistical Thinking for the 21st Century.*



Wasserman, L. (2013).  
*All of statistics: a concise course in statistical inference.*  
Springer Science & Business Media.